

Cortex Feature List

The private AI control plane for real business work

Cortex gives your business a governed AI workspace: one place to route requests, control models, manage context, supervise agents, approve actions, and turn scattered work into repeatable systems.

1. Proprietary Context Management

Cortex is built to control what the model sees, when it sees it, and how much context is actually worth spending.

Instead of dumping entire files, transcripts, histories, and tool outputs into every prompt, Cortex uses structured context handling, summaries, cached reads, tool-output shaping, session memory, and retrieval surfaces to keep AI work focused.

Why it matters: lower token waste, lower cost, faster runs, cleaner prompts, and less unnecessary exposure of sensitive information.

2. Intent-Aware Routing

Cortex does not treat every request the same.

An intent layer evaluates what the user is actually trying to do — summarize, extract, classify, compare, draft, search, route, automate, approve, or execute — and then chooses the right path.

That path may use a local model, a frontier model, a private cloud model, a tool, a workflow, a human approval gate, or a delegated agent.

Why it matters: better outcomes without asking employees to understand model selection, token budgets, tool routing, or infrastructure tradeoffs.

3. Model-Independent Execution

Cortex is not tied to one AI provider.

It can work across local models, open-weight models, private deployments, and frontier APIs. The business keeps the control layer even as models, prices, policies, and providers change.

Why it matters: no single-vendor dependency, better cost control, and the flexibility to use the right model for the right job.

4. Human-in-the-Loop Governance

Cortex is designed for supervised automation.

Approvals, audit logs, transcripts, budgets, routing rules, and execution boundaries make it clear what happened, what was approved, and where work went.

Why it matters: AI can move faster without removing human judgment from sensitive or high-value decisions.

5. Custody-First Security

Cortex keeps business control at the center.

Credentials, files, prompts, workflows, browser sessions, routing rules, and execution paths can remain under the client's environment and governance model. Sensitive information is handled through controlled access patterns instead of casual prompt exposure.

Why it matters: the model cannot leak what it never received.

6. Fling: Capture Work From Anywhere

Fling is Cortex's proprietary capture-and-routing layer.

An operator can grab useful context from a webpage, feed card, document, browser session, or selected content and send it to the right destination: intake, chat, job queue, webhook, feed, agent, or workflow.

Why it matters: useful work does not stay trapped in browser tabs, screenshots, links, email threads, and manual copy-paste.

7. Signal-to-Work Intake

Cortex can turn external signals into actionable work.

Browser captures, pasted text, links, social traffic, email, feed items, documents, and media can pass through trust checks, routing rules, approval gates, and landing queues before becoming tasks or automations.

Why it matters: Cortex does not just answer questions. It receives operational signal and helps route it into useful work.

8. Ranked Feeds and Work Queues

Cortex includes feed and queue surfaces that help operators review, prioritize, and act on incoming items.

Feeds can combine daily briefs, captured content, intake candidates, media, social items, and work-ready signals into a governed review flow.

Why it matters: teams get an operating inbox for AI-assisted work instead of another pile of disconnected notifications.

9. Multi-Tool Agent Workspace

Cortex supports serious tool-using agents.

Agents can use first-party tools, run scripts, search workspaces, read and send email, operate through APIs, enqueue jobs, and execute structured workflows under supervision.

Why it matters: Cortex moves beyond chat into governed execution.

10. Workflow Orchestration

Cortex can break complex work into phases: route, plan, execute, critique, archive, and improve.

Different models or tools can be used for different phases, with approvals and checkpoints where needed.

Why it matters: complex business work becomes repeatable instead of dependent on one person manually coordinating everything.

11. Local + Remote Deployment Flexibility

Cortex can operate locally, remotely, or in hybrid patterns depending on the client's security, cost, performance, and integration needs.

Why it matters: businesses can adopt AI without reorganizing their entire workflow around a hosted black box.

12. Workspace and Filemap Awareness

Cortex can help organize and reason across business files, project structures, documents, and working context.

Instead of forcing users to manually locate and paste everything, Cortex can build a more usable map of the work environment.

Why it matters: hidden knowledge becomes available at the moment of need.

13. Email, Browser, Social, and Media Integrations

Cortex is built to connect with the places work already happens: email, browser sessions, social channels, files, feeds, media workflows, job queues, and scriptable APIs.

Why it matters: automation fits the business instead of forcing the business into one app.

14. Team Sessions and Delegated Workers

Cortex can coordinate multiple controlled work sessions or delegated workers, each with its own task, boundary, and trust level.

Why it matters: larger tasks can be split across agents without losing oversight.

15. Script SDK, API, and MCP Surface

Cortex exposes developer-facing surfaces for scripts, extensions, tools, and host applications.

That means Cortex can become a shared execution layer underneath custom workflows, internal tools, and future Carapace interfaces.

Why it matters: the platform can grow with the client instead of staying trapped as a single chat window.

16. Auditability and Continuous Improvement

Cortex is built around inspectable work: transcripts, logs, routing decisions, approvals, feedback, and roadmap-driven improvements.

Why it matters: AI operations become something the business can review, govern, tune, and trust over time.

Suggested Feature Section Headline

Intelligence you control. Workflows you can trust.

Cortex is not another AI chat product. It is a private control layer for AI-assisted work: context-aware, model-independent, approval-gated, and designed around custody from the start.

Suggested Short Feature Grid

- Proprietary context management
- Intent-aware model routing
- Local and frontier model support

- Human approval gates
- Credential-conscious execution
- Audit logs and transcripts
- Fling browser capture
- Intake-to-work routing
- Ranked feeds and queues
- Multi-tool agents
- Workflow orchestration
- Email, browser, social, and media integrations
- Team sessions and delegated workers
- Script SDK, API, and MCP support
- Workspace and filemap awareness
- Continuous tuning and governance